

Neural networks

Victor Kitov

v.v.kitov@yandex.ru

Yandex School of Data Analysis



Table of Contents

- 1 Introduction
- 2 Output generation
- 3 Neural network optimization
- 4 Backpropagation algorithm
- 5 Invariances
- 6 Dropout

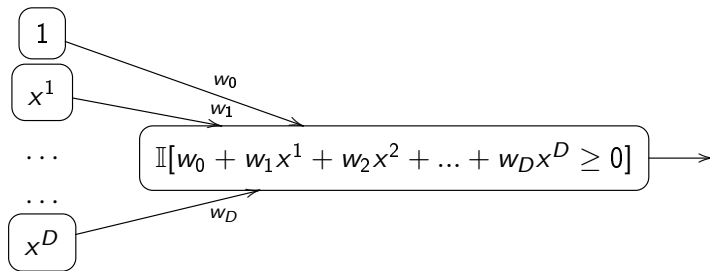
History

- Neural networks originally appeared as an attempt to model human brain



- Human brain consists of multiple interconnected neuron cells
 - cerebral cortex (the largest part) is estimated to contain 15–33 billion neurons
 - communication is performed by sending electrical and electro-chemical signals
 - signals are transmitted through axons - long thin parts of neurons.

Simple model of a neuron

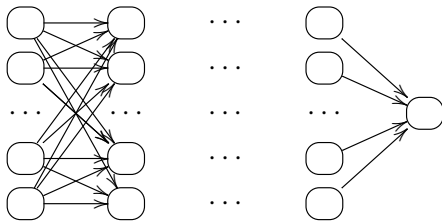


- Neuron get's activated in the half-space, defined by $w_0 + w_1 x^1 + w_2 x^2 + \dots + w_D x^D \geq 0$.
- Each node is called a neuron.
- Each edge is associated a weight.
- Constant feature 1 stands for bias.

Multilayer perceptron architecture¹

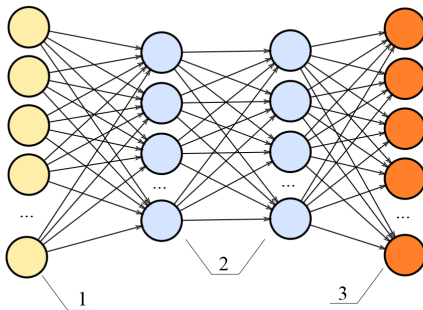
Multilayer perceptron:

- Each node has its own weights.
- Nodes&edges form directed acyclic graph.
- Multiple layers of neurons.
- Neurons of two consecutive layers are fully connected.



¹Propose neural networks estimating OR,AND,XOR functions on boolean inputs.

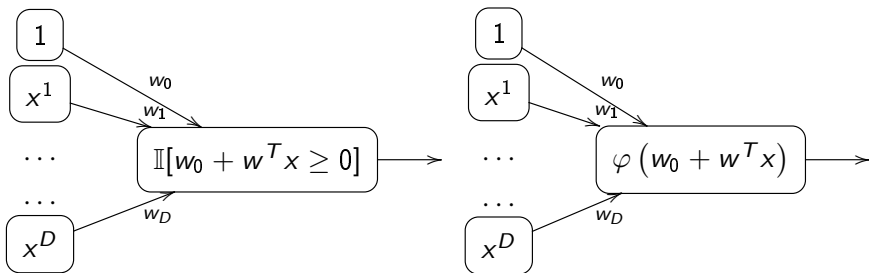
Layers



- Structure of neural network:
 - 1-input layer
 - 2-hidden layers
 - 3-output layer

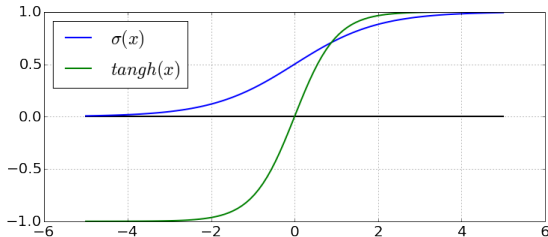
Continuous activations

- Pitfall of $\mathbb{I}[\cdot]$: it causes stepwise constant outputs, weight optimization methods become inapplicable.
- We can replace $\mathbb{I}[w^T x + w_0 \geq 0]$ with smooth activation $\varphi(w^T x + w_0)$



Typical activation functions

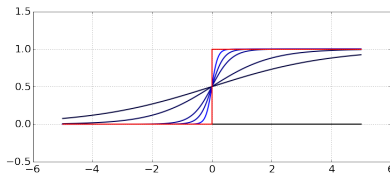
- sigmoidal: $\sigma(x) = \frac{1}{1+e^{-x}}$
 - 1-layer neural network with sigmoidal activation is equivalent to logistic regression
- hyperbolic tangent: $\tanh(x) = \frac{e^x - e^{-x}}{e^x + e^{-x}}$



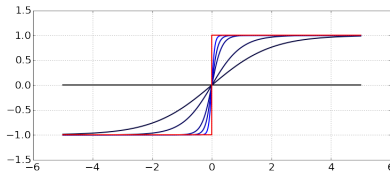
- ReLu: $\varphi(x) = [x]_+$.

Activation functions

Activation functions are smooth approximations of step functions:



$\sigma(ax)$ limits to 0/1-step function as $a \rightarrow \infty$



$\tanh(ax)$ limits to -1/1-step function as $a \rightarrow \infty$

Definition details

- Label each neuron with integer j .
- Denote: I_j - input to neuron j , O_j - output of neuron j
- Output of neuron j : $O_j = \varphi(I_j)$.
- Input to neuron j : $I_j = \sum_{k \in \text{inc}(j)} w_{kj} O_k + w_{0j}$,
 - w_{j0} is the bias term
 - $\text{inc}(j)$ is a set of neurons with outgoing edges incoming to neuron j .
 - further we will assume that at each layer there is a vertex with constant output $O_{\text{const}} \equiv 1$, so we can simplify notation

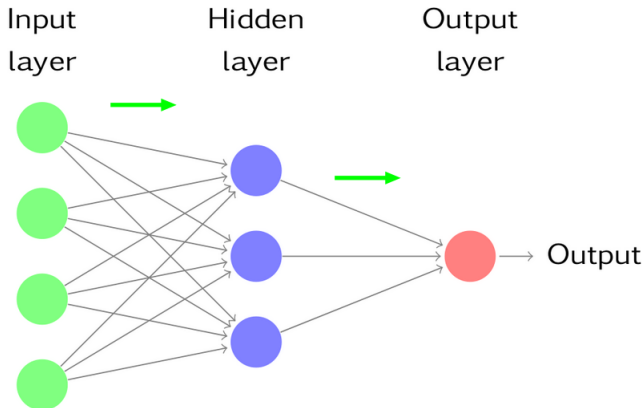
$$I_j = \sum_{k \in \text{inc}(j)} w_{kj} O_k$$

Table of Contents

- 1 Introduction
- 2 Output generation
- 3 Neural network optimization
- 4 Backpropagation algorithm
- 5 Invariances
- 6 Dropout

Output generation

- Forward propagation is a process of successive calculations of neuron outputs for given features.



Activations at output layer

- Regression: $\varphi(I) = I$
- Classification:
 - binary: $y \in \{+1, -1\}$

$$\varphi(I) = p(y = +1|x) = \frac{1}{1 + e^{-I}}$$

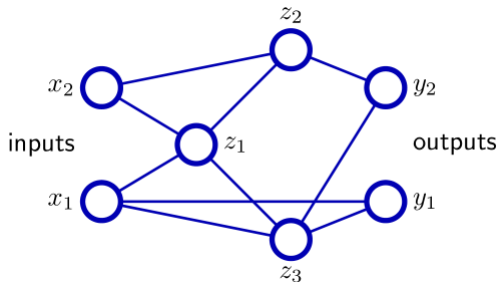
- multiclass: $y \in \{1, 2, \dots, C\}$

$$\varphi(O_1, \dots, O_C) = p(y = j|x) = \frac{e^{O_j}}{\sum_{k=1}^C e^{O_k}}, j = 1, 2, \dots, C$$

where O_1, \dots, O_C are outputs of output layer.

Generalizations

- each neuron j may have custom non-linear transformation φ_j
- weights may be constrained:
 - non-negative
 - equal weights
 - etc.
- layer skips are possible



- recurrent networks, RBF-activations.

Number of layers selection

- Number of layers usually denotes all layers except input layer (hidden layers+output layer)
- Classification with indicator activations:
 - single layer network selects arbitrary half-spaces
 - 2-layer network selects arbitrary convex polyhedron (by intersection of 1-layer outputs)
 - therefore it can approximate arbitrary convex sets
 - 3-layer network selects (by union of 2-layer outputs) arbitrary finite sets of polyhedra
 - therefore it can approximate all measurable sets

Number of layers selection

- Regression with indicator activations:
 - single layer can approximate arbitrary linear function
 - 2-layer network can model indicator function of arbitrary convex polyhedron
 - 3-layer network can uniformly approximate arbitrary continuous function (as sum weighted sum of indicators convex polyhedra)

Sufficient amount of layers

Any continuous function on a compact space can be uniformly approximated by 2-layer neural network with linear output and wide range of activation functions (excluding polynomial).

- In practice often it is more convenient to use more layers with less total amount of neurons
 - model becomes more interpretable and easy to fit.

Table of Contents

- 1 Introduction
- 2 Output generation
- 3 Neural network optimization**
- 4 Backpropagation algorithm
- 5 Invariances
- 6 Dropout

Network optimization: regression

- Single output:

$$\frac{1}{N} \sum_{n=1}^N (\hat{y}_n(x_n|w) - y_n)^2 \rightarrow \min_w$$

Network optimization: regression

- Single output:

$$\frac{1}{N} \sum_{n=1}^N (\hat{y}_n(x_n|w) - y_n)^2 \rightarrow \min_w$$

- K outputs

$$\frac{1}{NK} \sum_{n=1}^N \sum_{k=1}^K (\hat{y}_{nk}(x_n|w) - y_{nk})^2 \rightarrow \min_w$$

Network optimization: classification

- Two classes, $y \in \{0, 1\}$:

$$\prod_{n=1}^N p(y_n = 1|x_n, w)^{y_n} [1 - p(y_n = 1|x_n, w)]^{1-y_n} \rightarrow \max_w$$

Network optimization: classification

- Two classes, $y \in \{0, 1\}$:

$$\prod_{n=1}^N p(y_n = 1|x_n, w)^{y_n} [1 - p(y_n = 1|x_n, w)]^{1-y_n} \rightarrow \max_w$$

- C classes, $y_{nc} = \mathbb{I}\{y_n = c\}$:

$$\prod_{n=1}^N \prod_{c=1}^C p(y_n = c|x_n, w)^{y_{nc}} \rightarrow \max_w$$

Network optimization: classification

- Two classes, $y \in \{0, 1\}$:

$$\prod_{n=1}^N p(y_n = 1|x_n, w)^{y_n} [1 - p(y_n = 1|x_n, w)]^{1-y_n} \rightarrow \max_w$$

- C classes, $y_{nc} = \mathbb{I}\{y_n = c\}$:

$$\prod_{n=1}^N \prod_{c=1}^C p(y_n = c|x_n, w)^{y_{nc}} \rightarrow \max_w$$

- In practice log-likelihood is maximized to avoid numerical underflow.

Neural network optimization

- Denote W - the total dimensionality of weights space, $w \in \mathbb{R}^W$.
- Let $E(\hat{y}, y)$ denote the loss function of output
- We may optimize neural network using gradient descent:

```
k = 0
initialize randomly  $w^0$  # small values for sigmoid and tanh

while (stop criteria not met):
     $w^{k+1} := w^k - \eta \nabla E(w^k)$ 
    k := k + 1
```

- Standardization of features makes gradient descend converge faster
- Other optimization methods are more efficient (such as conjugate gradients)

Gradient calculation

- Denote $\varepsilon_i = (0, \dots, 0, \overbrace{\varepsilon}^{i\text{-th position}}, 0, \dots, 0) \in \mathbb{R}^W$
- Direct $\nabla E(w)$ calculation, using

$$\frac{\partial E}{\partial w_i} = \frac{E(w + \varepsilon_i) - E(w)}{\varepsilon} + O(\varepsilon)$$

or better

$$\frac{\partial E}{\partial w_i} = \frac{E(w + \varepsilon_i) - E(w - \varepsilon_i)}{2\varepsilon} + O(\varepsilon^2)$$

has complexity $O(W^2)$ [W forward propagations to evaluate W derivatives]

Backpropagation algorithm needs only $O(W)$ to evaluate all derivatives.

Multiple local optima problem

- Optimization problem for neural nets is **non-convex**.
- Different optima will correspond to:
 - different starting parameter values
 - different training samples
- So we may solve task many times for different conditions and then
 - select best model
 - alternatively: average different obtained models to get ensemble

Table of Contents

- 1 Introduction
- 2 Output generation
- 3 Neural network optimization
- 4 Backpropagation algorithm**
- 5 Invariances
- 6 Dropout

Definitions

- Denote w_{ij} be the weight of edge, connecting i -th and j -th neuron.
- Define $\delta_j = \frac{\partial E}{\partial I_j} = \frac{\partial E}{\partial O_j} \frac{\partial O_j}{\partial I_j}$
- Since E depends on w_{ij} through the following functional relationship $E(w_{ij}) \equiv E(O_j(I_j(w_{ij})))$, using the chain rule we obtain:

$$\frac{\partial E}{\partial w_{ij}} = \frac{\partial E}{\partial I_j} \frac{\partial I_j}{\partial w_{ij}} = \delta_j O_i$$

because $\frac{\partial I_j}{\partial w_{ij}} = \frac{\partial}{\partial w_{ij}} \left(\sum_{k \in \text{inc}(j)} w_{kj} O_k \right) = O_i$, where $\text{inc}(j)$ is a set of all neurons with outgoing edges to neuron j .

- $\frac{\partial E}{\partial I_j} = \frac{\partial E}{\partial O_j} \frac{\partial O_j}{\partial I_j} = \frac{\partial E}{\partial O_j} \varphi'(I_j)$, where φ is the activation function.

Output layer

- If neuron j belongs to the output node, then error $\frac{\partial E}{\partial O_j}$ is calculated directly.
- For output layer deltas are calculated directly:

$$\delta_j = \frac{\partial E}{\partial O_j} \frac{\partial O_j}{\partial I_j} = \frac{\partial E}{\partial O_j} \varphi'(I_j) \quad (1)$$

- example for training set = {single point x and true vector of outputs $(y_1, \dots, y_{|O_L|})$ }:
 - for $E = \frac{1}{2} \sum_{j \in O_L} (O_j - y_j)^2$:

$$\frac{\partial E}{\partial O_j} = O_j - y_j$$

- for $\varphi(I) = \text{sigm}(I)$:

$$\varphi'(I_j) = \sigma(I_j) (1 - \sigma(I_j)) = O_j(1 - O_j)$$

- finally

$$\delta_j = (O_j - y_j) O_j(1 - O_j)$$

Inner layer

- If neuron j belongs some hidden layer, denote $out(j) = \{k_1, k_2, \dots, k_m\}$ the set of all neurons, receiving output from neuron j .
- The effect of O_j on E is fully absorbed by $l_{k_1}, l_{k_2}, \dots, l_{k_m}$, so

$$\frac{\partial E(O_j)}{\partial O_j} = \frac{\partial E(l_{k_1}, l_{k_2}, \dots, l_{k_m})}{\partial O_j} = \sum_{k \in out(j)} \left(\frac{\partial E}{\partial l_k} \frac{\partial l_k}{\partial O_j} \right) = \sum_{k \in out(j)} (\delta_k w_{jk})$$

- So for layers other than output layer we have:

$$\delta_j = \frac{\partial E}{\partial l_j} = \frac{\partial E}{\partial O_j} \frac{\partial O_j}{\partial l_j} = \sum_{k \in out(j)} (\delta_k w_{jk}) \varphi'(l_j) \quad (2)$$

- Weight derivatives are calculated using errors and outputs:

$$\frac{\partial E}{\partial w_{ij}} = \frac{\partial E}{\partial l_j} \frac{\partial l_j}{\partial w_{ij}} = \delta_j O_i \quad (3)$$

Backpropagation

- Backpropagation algorithm:
 - 1 Forward propagate x_n to the neural network, store all inputs I_i and outputs O_i for each neuron.
 - 2 Calculate δ_i for all $i \in \text{output layer}$ using (1).
 - 3 Backpropagate δ_i from final layer backwards layer by layer using (2).
 - 4 Using calculated deltas and outputs calculate $\frac{\partial E}{\partial w_{ij}}$ with (3).
- Algorithm complexity: $O(W)$, where W is total number of edges.
- Updates:
 - batch
 - stochastic
 - using minibatches of objects

Regularization

- Constrain model complexity directly
 - constrain number of neurons
 - constrain number of layers
 - impose constraints on weights
- Take a flexible model
 - use early stopping during iterative evaluation (by controlling validation error)
 - quadratic regularization

$$\tilde{E}(w) = E(w) + \lambda \sum_i w_i^2$$

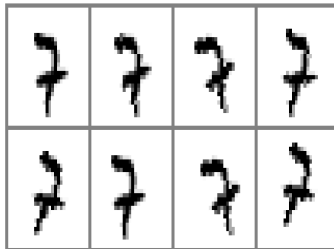
- Stochastic simplification of model, such as dropout.

Table of Contents

- 1 Introduction
- 2 Output generation
- 3 Neural network optimization
- 4 Backpropagation algorithm
- 5 Invariances**
- 6 Dropout

Invariance to particular transformations

- It may happen that solution should not depend on certain kinds of transformations in the input space.
- Example: character recognition task
 - translation invariance
 - scale invariance
 - invariance to small rotations
 - invariance to small uniform noise



Invariance to particular transformations

- Approaches to build an invariant model to particular transformation:
 - use only features that are invariant to transformations
 - augment training objects with their transformed copies according to given invariances
 - amount of possible transformations grows exponentially with the number of invariances
 - add regularization term to the target cost function, which penalizes changes in output after invariant transformations
 - see tangent propagation
 - build the invariance properties into the structure of neural network
 - see convolutional neural networks

Augmentation of training samples

- 1 generate a random set of invariant transformations
- 2 apply these transformations to training objects
- 3 obtain new training objects

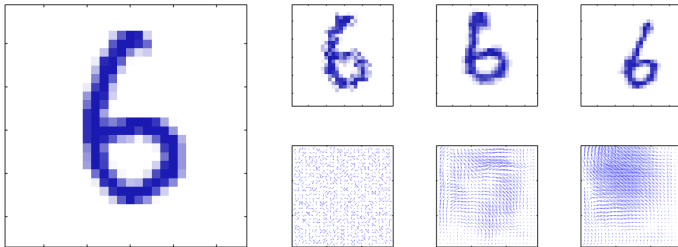
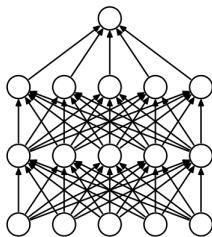


Table of Contents

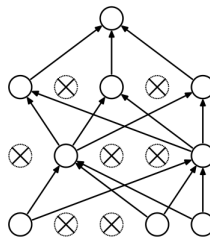
- 1 Introduction
- 2 Output generation
- 3 Neural network optimization
- 4 Backpropagation algorithm
- 5 Invariances
- 6 Dropout**

Dropout idea

Each neuron is left or removed independently from decisions about other nodes:



(a) Standard Neural Net



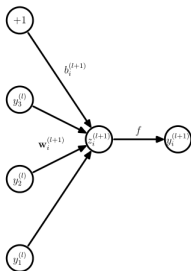
(b) After applying dropout.

- Output layer nodes are never removed.
- Recommended² neuron removal probability $1 - p$:
 - $p = 0.5$ for inner layer nodes
 - $p = 0.8$ for input layer nodes (feature subsampling)

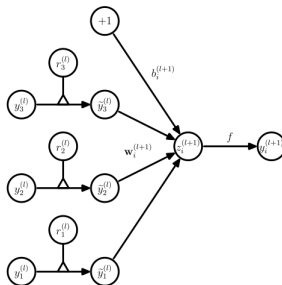
²Removal probabilities can be fine-tuned on cross-validation.

Dropout algorithm

- Dropout forces neurons to rely less on other neurons and to learn something valuable by themselves.
 - neurons learning becomes more robust to information learnt by other neurons.
 - resulting network becomes less overfitted.
- $r_i^{(l)}$ are i.i.d. random variables having *Bernoulli*(p) distribution.



(a) Standard network



(b) Dropout network

Definitions

Define:

- $f(x)$ - an activation function.
- y^l - vector of outputs at layer l
- z^l - vector of inputs to layer l
- L - number of layers in neural network
- $y^{(0)} = x$ - input feature vector
- $Bernoulli(p)$ returns a vector of independent Bernoulli random variables with parameter p .

Forward propagation algorithm

We need to repeat forward propagation recurrently for $l = 0, 1, \dots, L - 1$.

- 1 Usual feed-forward neural network:

$$z_i^{(l+1)} = \left(w_i^{(l+1)} \right)^T y^l + b_i^{(l+1)}$$

$$y_i^{(l+1)} = f(z_i^{(l+1)})$$

- 2 Feed-forward network with dropout:

$$\mathbf{r}_j^{(l)} \sim \text{Bernoulli}(\mathbf{p})$$

$$\tilde{\mathbf{y}}^l = \mathbf{r}^{(l)} * \mathbf{y}^{(l)}$$

$$z_i^{(l+1)} = \left(w_i^{(l+1)} \right)^T \tilde{\mathbf{y}}^l + b_i^{(l+1)}$$

$$y_i^{(l+1)} = f(z_i^{(l+1)})$$

Application of dropout

- **Learning**

- while weights not converge:

- 1 sample random subnetwork (“thinned network”) with dropout
- 2 apply one step of stochastic gradient descent to thinned network

Comment: due to weights sharing across all thinned networks the number of parameters is the same as in original network.

- **Prediction**

- use full networks with all nodes, but multiply each weight by p
- such scaling approximates average output of randomly sampled thinned network³.

³Approximation is precise for networks without non-linearities. With non-linearities - not. We may also use Monte-Carlo sampling in the latter case.

Conclusion on neural networks

- Advantages of neural networks:
 - can model accurately complex non-linear relationships
 - easily parallelizable
- Disadvantages of neural networks:
 - hardly interpretable (“black-box” algorithm)
 - optimization requires skill
 - too many parameters
 - may converge slowly
 - may converge to inefficient local minimum far from global one